

CREATING A DATA MODEL

CREATING A DATA MODEL



In this section we'll cover **foundational data modeling topics** like normalization, fact and dimension tables, primary and foreign keys, relationship cardinality and filter flow

TOPICS WE'LL COVER:

Data Modeling 101

Normalization

Facts & Dimensions

Primary & Foreign Keys

Cardinality

Filter Flow

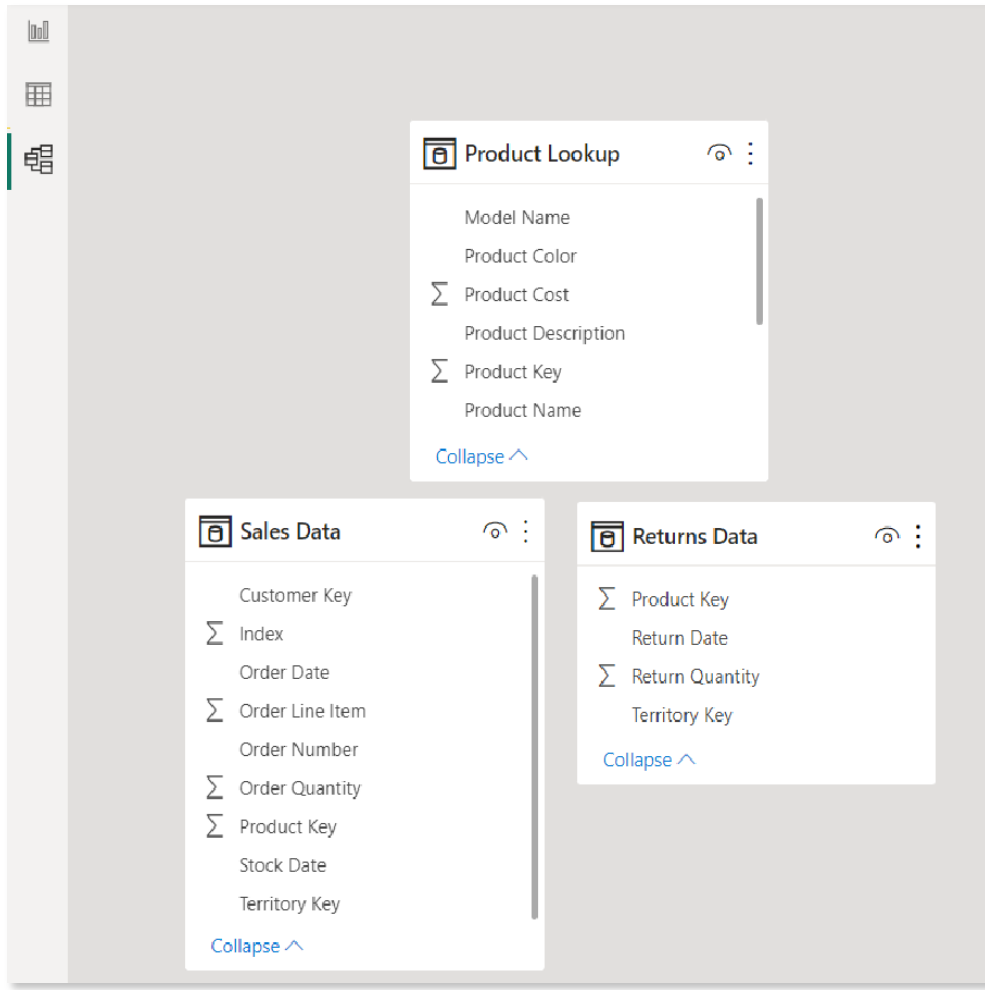
Common Schemas

Hierarchies

GOALS FOR THIS SECTION:

- Understand the basic principles of data modeling, including normalization, fact & dimension tables and common schemas
- Create table relationships using primary and foreign keys, and discuss different types of relationship cardinality
- Configure report filters and trace filter context as it flows between related tables in the model
- Explore data modeling options like hierarchies, data categories and hidden fields

WHAT IS A DATA MODEL?



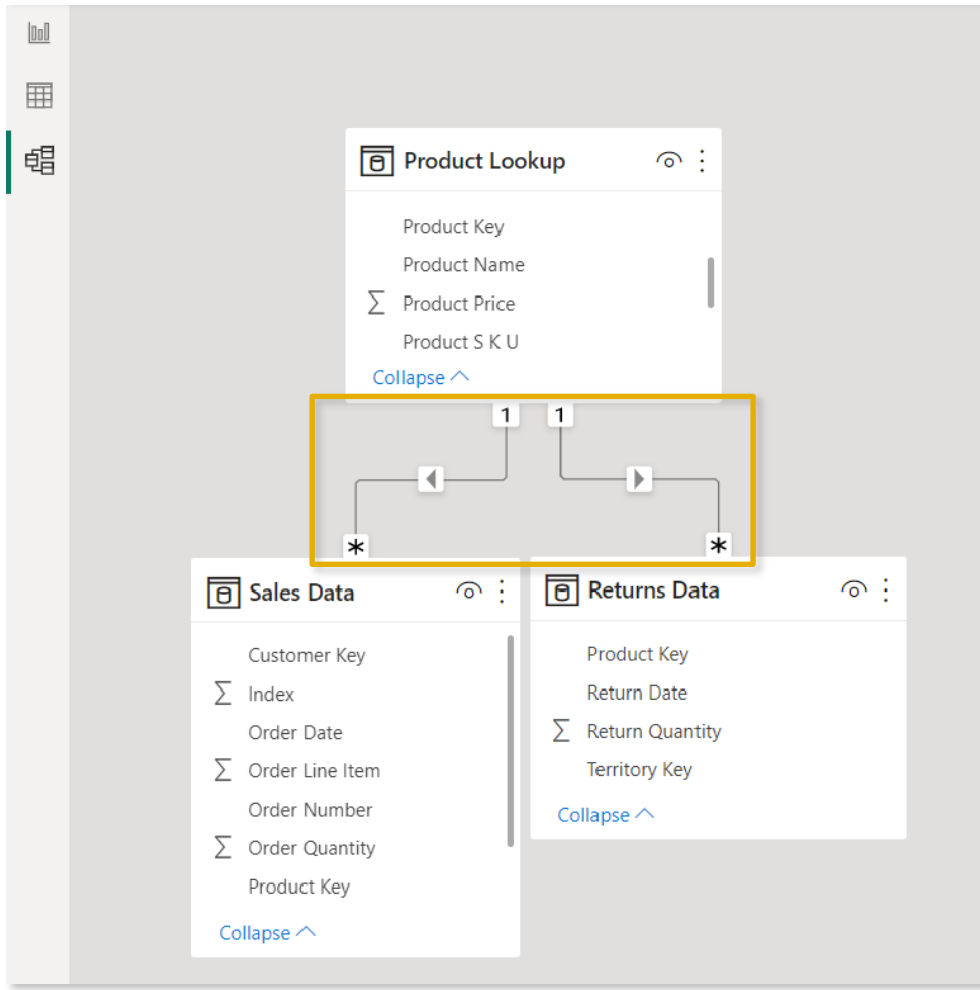
This **IS NOT** a data model



- This is a collection of independent tables, which share no connections or relationships
- If you tried to visualize **Orders** and **Returns** by **Product**, this is what you'd get

| ProductName | OrderQuantity | ReturnQuantity |
|------------------------|---------------|----------------|
| All-Purpose Bike Stand | 84,174 | 1,828 |
| AWC Logo Cap | 84,174 | 1,828 |
| Bike Wash - Dissolver | 84,174 | 1,828 |
| Cable Lock | 84,174 | 1,828 |
| Chain | 84,174 | 1,828 |
| Classic Vest, L | 84,174 | 1,828 |
| Classic Vest, M | 84,174 | 1,828 |
| Classic Vest, S | 84,174 | 1,828 |
| Fender Set - Mountain | 84,174 | 1,828 |
| Total | 84,174 | 1,828 |

WHAT IS A DATA MODEL?



This **IS** a data model! 😊

- The tables are connected via relationships, based on a common field (Product Key)
- Now **Sales** and **Returns** data can be filtered using fields from the **Product Lookup** table!

| ProductName | OrderQuantity | ReturnQuantity |
|------------------------|---------------|----------------|
| All-Purpose Bike Stand | 234 | 8 |
| AWC Logo Cap | 4,151 | 46 |
| Bike Wash - Dissolver | 1,706 | 25 |
| Classic Vest, L | 182 | 4 |
| Classic Vest, M | 182 | 7 |
| Classic Vest, S | 157 | 8 |
| Fender Set - Mountain | 3,960 | 54 |
| Half-Finger Gloves, L | 840 | 18 |
| Half-Finger Gloves, M | 918 | 16 |
| Total | 84,174 | 1,828 |

DATABASE NORMALIZATION



Normalization is the process of organizing the tables and columns in a relational database to reduce redundancy and preserve data integrity. It's commonly used to:

- **Eliminate redundant data** to decrease table sizes and improve processing speed & efficiency
- **Minimize errors and anomalies** from data modifications (inserting, updating or deleting records)
- **Simplify queries** and structure the database for meaningful analysis



In a normalized database, each table should serve a **distinct** and **specific** purpose (*i.e. product information, transaction records, customer attributes, store details, etc.*)

| date | product_id | quantity | product_brand | product_name | product_sku | product_weight |
|----------|------------|----------|---------------|-----------------------------|-------------|----------------|
| 1/1/1997 | 869 | 5 | Nationeel | Nationeel Grape Fruit Roll | 52382137179 | 17 |
| 1/7/1997 | 869 | 2 | Nationeel | Nationeel Grape Fruit Roll | 52382137179 | 17 |
| 1/3/1997 | 1 | 4 | Washington | Washington Berry Juice | 90748583674 | 8.39 |
| 1/1/1997 | 1472 | 3 | Fort West | Fort West Fudge Cookies | 37276054024 | 8.28 |
| 1/6/1997 | 1472 | 2 | Fort West | Fort West Fudge Cookies | 37276054024 | 8.28 |
| 1/5/1997 | 2 | 4 | Washington | Washington Mango Drink | 96516502499 | 7.42 |
| 1/1/1997 | 76 | 4 | Red Spade | Red Spade Sliced Chicken | 62054644227 | 18.1 |
| 1/1/1997 | 76 | 2 | Red Spade | Red Spade Sliced Chicken | 62054644227 | 18.1 |
| 1/5/1997 | 3 | 2 | Washington | Washington Strawberry Drink | 58427771925 | 13.1 |
| 1/7/1997 | 3 | 2 | Washington | Washington Strawberry Drink | 58427771925 | 13.1 |
| 1/1/1997 | 320 | 3 | Excellent | Excellent Cranberry Juice | 36570182442 | 16.4 |

Models that aren't normalized contain **redundant, duplicate data**. In this case, all of the product-specific fields could be stored in a separate table containing a unique record for each **product id**

This may not seem critical now, but minor inefficiencies can become major problems at scale!



FACT & DIMENSION TABLES

Data models generally contain two types of tables: **fact** (“data”) tables, and **dimension** (“lookup”) tables:

- **Fact tables** contain **numerical values** or metrics used for summarization (*sales, orders, transactions, pageviews, etc.*)
- **Dimension tables** contain **descriptive attributes** used for filtering or grouping (*products, customers, dates, stores, etc.*)

| date | product_id | quantity |
|----------|------------|----------|
| 1/1/1997 | 869 | 5 |
| 1/1/1997 | 1472 | 3 |
| 1/1/1997 | 76 | 4 |
| 1/1/1997 | 320 | 3 |
| 1/1/1997 | 4 | 4 |
| 1/1/1997 | 952 | 4 |
| 1/1/1997 | 1222 | 4 |
| 1/1/1997 | 517 | 4 |
| 1/1/1997 | 1359 | 4 |
| 1/1/1997 | 357 | 4 |
| 1/1/1997 | 1426 | 5 |
| 1/1/1997 | 190 | 4 |
| 1/1/1997 | 367 | 4 |
| 1/1/1997 | 250 | 5 |
| 1/1/1997 | 600 | 4 |
| 1/1/1997 | 702 | 5 |

This **Fact** table contains **quantity** values, along with **date** and **product_id** fields

| date | day_of_month | month | year | weekday | week_of_year | week_ending | month_name | quarter |
|----------|--------------|-------|------|-----------|--------------|-------------|------------|---------|
| 1/1/1997 | 1 | 1 | 1997 | Wednesday | 1 | 1/5/1997 | January | Q1 |
| 1/2/1997 | 2 | 1 | 1997 | Thursday | 1 | 1/5/1997 | January | Q1 |
| 1/3/1997 | 3 | 1 | 1997 | Friday | 1 | 1/5/1997 | January | Q1 |
| 1/4/1997 | 4 | 1 | 1997 | Saturday | 1 | 1/5/1997 | January | Q1 |
| 1/5/1997 | 5 | 1 | 1997 | Sunday | 2 | 1/5/1997 | January | Q1 |
| 1/6/1997 | 6 | 1 | 1997 | Monday | 2 | 1/12/1997 | January | Q1 |

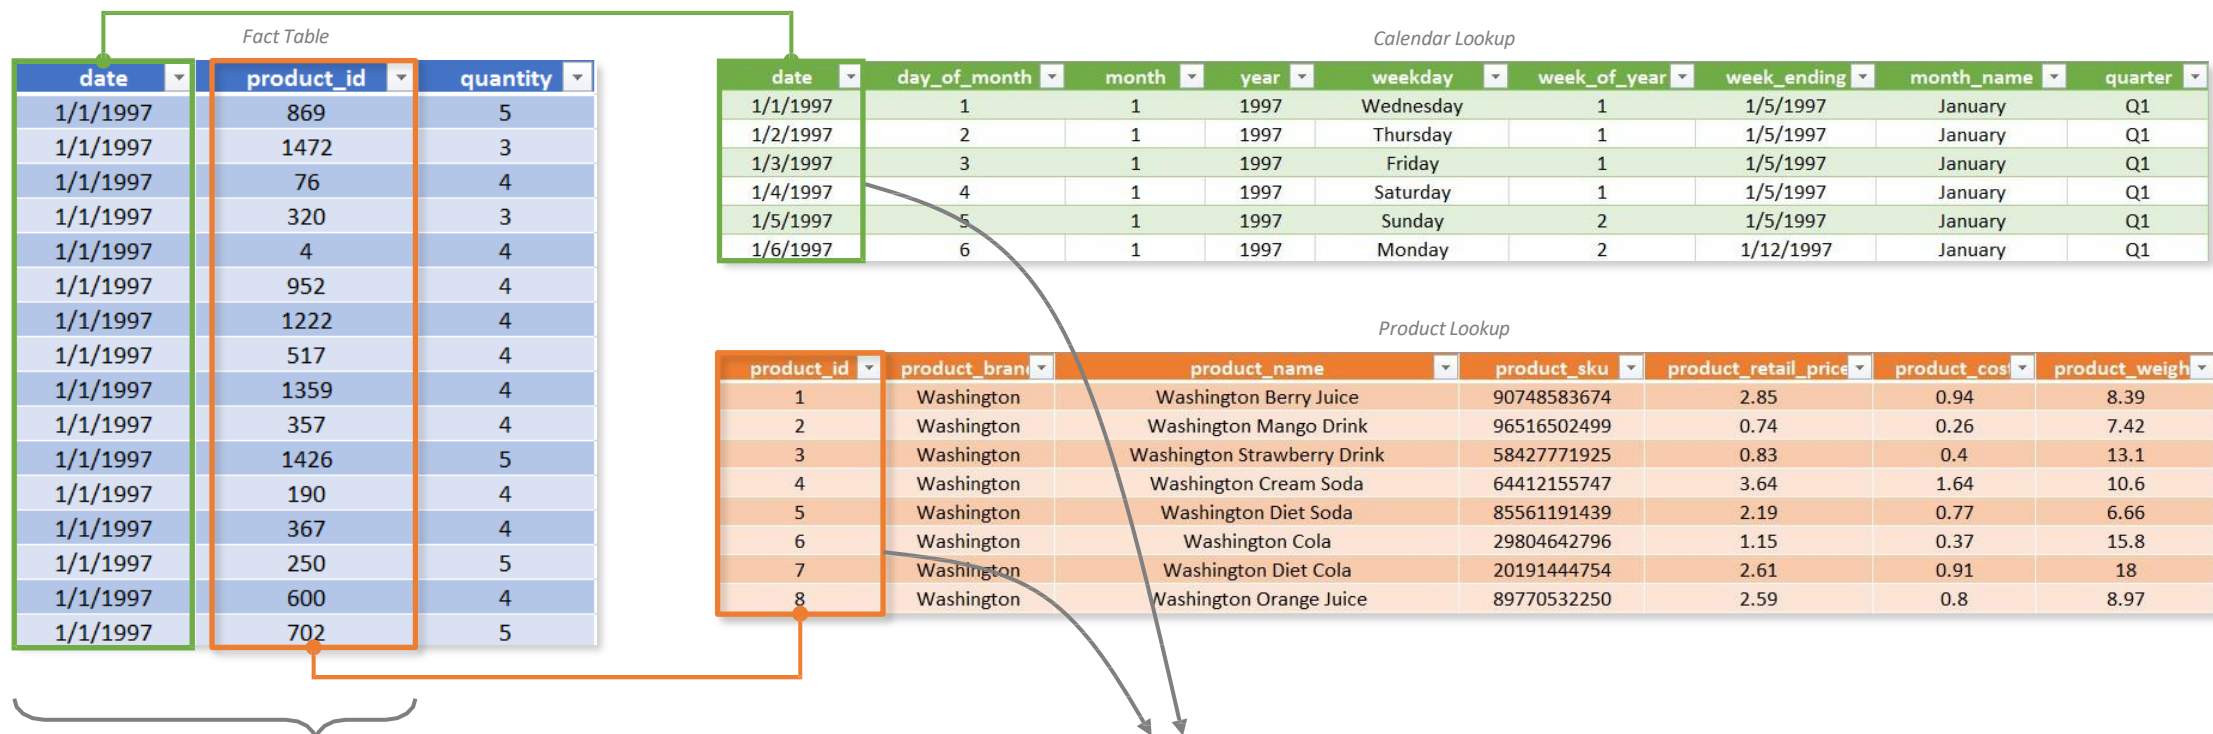
This **Calendar Lookup** table contains attributes about each **date** (month, year, quarter, etc.)

| product_id | product_brand | product_name | product_sku | product_retail_price | product_cost | product_weight |
|------------|---------------|-----------------------------|-------------|----------------------|--------------|----------------|
| 1 | Washington | Washington Berry Juice | 90748583674 | 2.85 | 0.94 | 8.39 |
| 2 | Washington | Washington Mango Drink | 96516502499 | 0.74 | 0.26 | 7.42 |
| 3 | Washington | Washington Strawberry Drink | 58427771925 | 0.83 | 0.4 | 13.1 |
| 4 | Washington | Washington Cream Soda | 64412155747 | 3.64 | 1.64 | 10.6 |
| 5 | Washington | Washington Diet Soda | 85561191439 | 2.19 | 0.77 | 6.66 |
| 6 | Washington | Washington Cola | 29804642796 | 1.15 | 0.37 | 15.8 |
| 7 | Washington | Washington Diet Cola | 20191444754 | 2.61 | 0.91 | 18 |
| 8 | Washington | Washington Orange Juice | 89770532250 | 2.59 | 0.8 | 8.97 |

This **Product Lookup** table contains attributes about each **product_id** (brand, SKU, price, etc.)



PRIMARY & FOREIGN KEYS



These are **foreign keys** (FK)

They contain multiple instances of each value, and relate to **primary keys** in dimension tables

These are **primary keys** (PK)

They uniquely identify each row of the table, and relate to **foreign keys** in fact tables

RELATIONSHIPS VS. MERGED TABLES



Can't I just merge queries or use lookup functions to **pull everything into one single table**?

- Anonymous confused man

Original **Fact Table** fields

Attributes from **Calendar Lookup** table

Attributes from **Product Lookup** table

| date | product_id | quantity | day_of_month | month | year | weekday | month_name | quarter | product_brand | product_name | product_sku | product_weight |
|----------|------------|----------|--------------|-------|------|-----------|------------|---------|---------------|-----------------------------|-------------|----------------|
| 1/1/1997 | 869 | 5 | 1 | 1 | 1997 | Wednesday | January | Q1 | Nationeel | Nationeel Grape Fruit Roll | 52382137179 | 17 |
| 1/7/1997 | 869 | 2 | 7 | 1 | 1997 | Tuesday | January | Q1 | Nationeel | Nationeel Grape Fruit Roll | 52382137179 | 17 |
| 1/3/1997 | 1 | 4 | 3 | 1 | 1997 | Friday | January | Q1 | Washington | Washington Berry Juice | 90748583674 | 8.39 |
| 1/1/1997 | 1472 | 3 | 1 | 1 | 1997 | Wednesday | January | Q1 | Fort West | Fort West Fudge Cookies | 37276054024 | 8.28 |
| 1/6/1997 | 1472 | 2 | 6 | 1 | 1997 | Monday | January | Q1 | Fort West | Fort West Fudge Cookies | 37276054024 | 8.28 |
| 1/5/1997 | 2 | 4 | 5 | 1 | 1997 | Sunday | January | Q1 | Washington | Washington Mango Drink | 96516502499 | 7.42 |
| 1/1/1997 | 76 | 4 | 1 | 1 | 1997 | Wednesday | January | Q1 | Red Spade | Red Spade Sliced Chicken | 62054644227 | 18.1 |
| 1/1/1997 | 76 | 2 | 1 | 1 | 1997 | Wednesday | January | Q1 | Red Spade | Red Spade Sliced Chicken | 62054644227 | 18.1 |
| 1/5/1997 | 3 | 2 | 5 | 1 | 1997 | Sunday | January | Q1 | Washington | Washington Strawberry Drink | 58427771925 | 13.1 |
| 1/7/1997 | 3 | 2 | 7 | 1 | 1997 | Tuesday | January | Q1 | Washington | Washington Strawberry Drink | 58427771925 | 13.1 |
| 1/1/1997 | 320 | 3 | 1 | 1 | 1997 | Wednesday | January | Q1 | Excellent | Excellent Cranberry Juice | 36570182442 | 16.4 |

You can, **but it's extremely inefficient!**

- Merging tables creates **redundancy** and often requires **significantly more memory and processing power** to analyze compared to a relational model with multiple small tables



THE MODEL VIEW

The screenshot displays the 'Model View' interface, which is used for visualizing and managing data models. The interface is divided into several key sections:

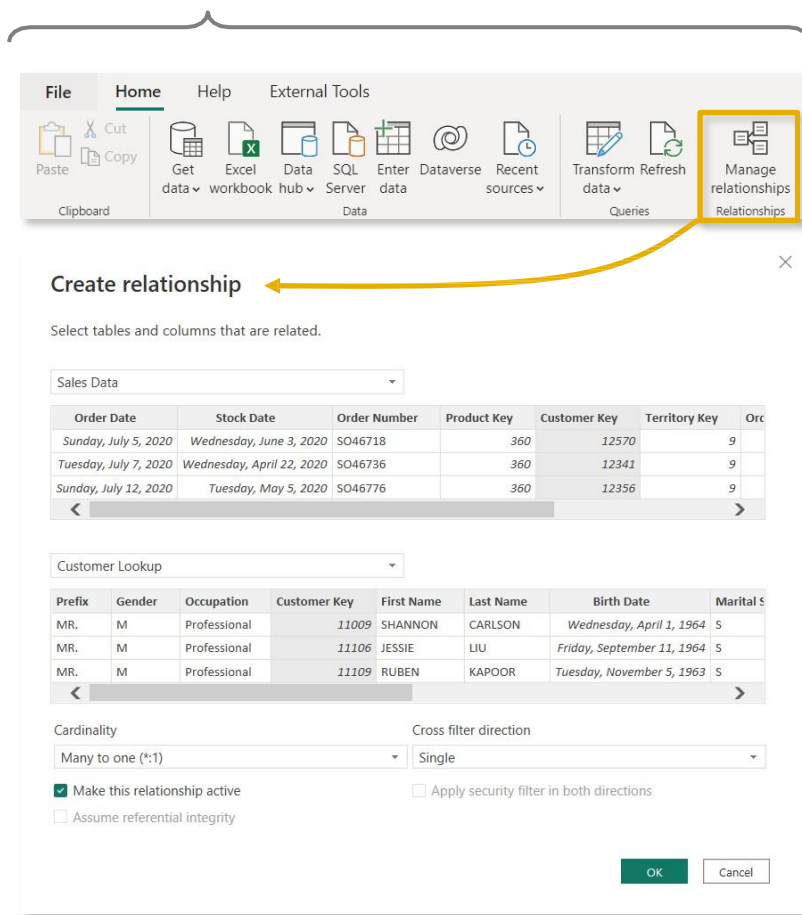
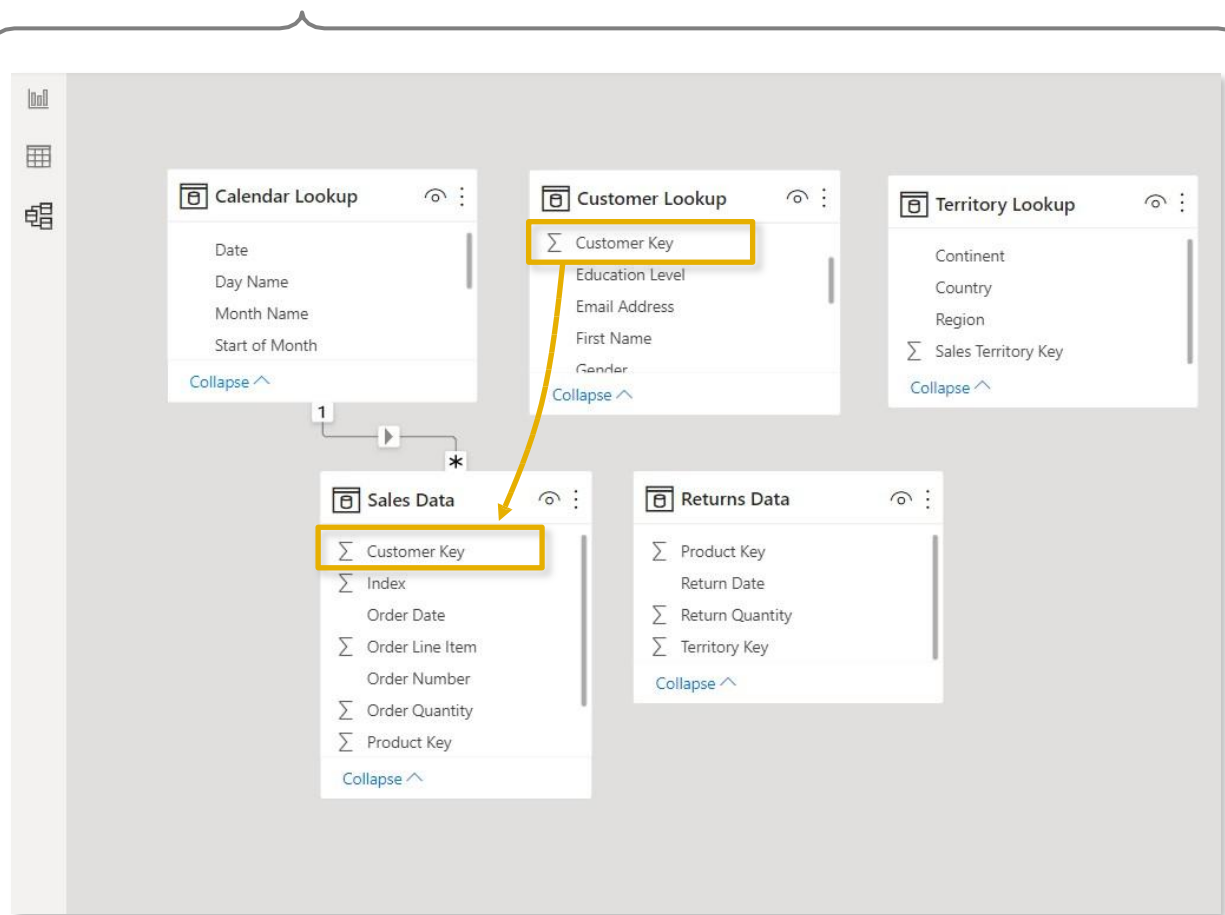
- Menu Ribbon (Home, Help):** Located at the top, this ribbon contains various toolbars for data management, queries, relationships, calculations, security, Q&A, and publishing.
- Model canvas:** The central workspace where the data model is visualized. It shows a hierarchy of tables: 'Calendar Lookup', 'Customer Lookup', 'Territory Lookup', 'Product Lookup', 'Sales Data', and 'Returns Data'. Lines connect these tables to represent relationships.
- Properties pane:** Located on the right side, this pane allows users to configure the selected table. It includes fields for Name, Description, Synonyms, Row label, Key column, and Is hidden.
- Data / Field List:** Also on the right, this pane provides a searchable list of fields available in the data source. Fields are categorized by table, and users can expand or collapse sections.
- Model layout tabs:** At the bottom left, there is a tab labeled 'All tables' with a plus sign, indicating where users can manage the layout of the model.
- View Options:** At the bottom right, there are controls for zooming in and out (indicated by a slider and '80%') and buttons for 'Reset Layout' and 'Fit to Page'.

CREATING TABLE RELATIONSHIPS



OPTION 1: Click and drag to connect primary and foreign keys within the **Model** view

OPTION 2: Add or detect relationships using the **Manage Relationships** dialog box



MANAGING & EDITING RELATIONSHIPS



Manage relationships

| Active | From: Table (Column) | To: Table (Column) |
|-------------------------------------|---|--|
| <input checked="" type="checkbox"/> | Product Lookup (Product Subcategory Key) | Product Subcategories Lookup (Product Subcategory Key) |
| <input checked="" type="checkbox"/> | Product Subcategories Lookup (Product Category Key) | Product Categories Lookup (Product Category Key) |
| <input checked="" type="checkbox"/> | Sales Data (Customer Key) | Customer Lookup (Customer Key) |
| <input checked="" type="checkbox"/> | Sales Data (Order Date) | Calendar Lookup (Date) |
| <input checked="" type="checkbox"/> | Sales Data (Product Key) | Product Lookup (Product Key) |

Buttons: New..., Autodetect..., **Edit...**, Delete

Close

Launch the **Manage Relationships** dialog box or double-click a relationship to modify it

Edit relationship

Select tables and columns that are related.

Sales Data

| Order Date | Stock Date | Order Number | Product Key | Customer Key | Territory Key | Order Date |
|-----------------------|---------------------------|--------------|-------------|--------------|---------------|------------|
| Sunday, July 5, 2020 | Wednesday, June 3, 2020 | SO46718 | 360 | 12570 | 9 | |
| Tuesday, July 7, 2020 | Wednesday, April 22, 2020 | SO46736 | 360 | 12341 | 9 | |
| Sunday, July 12, 2020 | Tuesday, May 5, 2020 | SO46776 | 360 | 12356 | 9 | |

Customer Lookup

| Prefix | Gender | Occupation | Customer Key | First Name | Last Name | Birth Date | Marital S |
|--------|--------|--------------|--------------|------------|-----------|----------------------------|-----------|
| MR. | M | Professional | 11009 | SHANNON | CARLSON | Wednesday, April 1, 1964 | S |
| MR. | M | Professional | 11106 | JESSIE | LIU | Friday, September 11, 1964 | S |
| MR. | M | Professional | 11109 | RUBEN | KAPOOR | Tuesday, November 5, 1963 | S |

Cardinality: Many to one (*:1) | Cross filter direction: Single

☒ Make this relationship active | ☐ Apply security filter in both directions
☐ Assume referential integrity

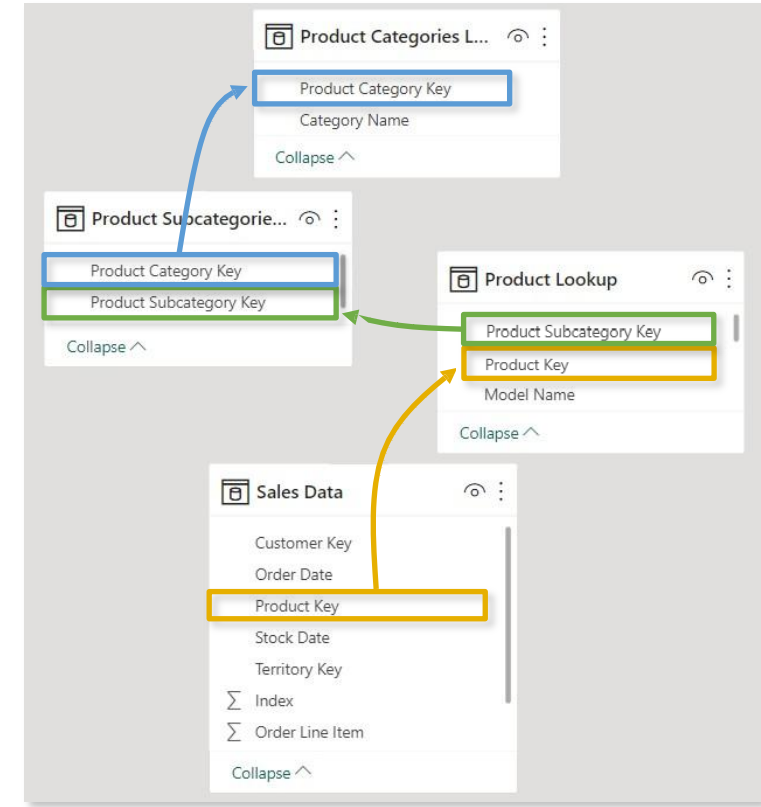
OK Cancel

Editing tools allow you to **activate or deactivate** relationships and manage **cardinality** and **filter direction** – more on that soon!

STAR & SNOWFLAKE SCHEMAS




A **star schema** is the simplest and most common type of data model, characterized by a single fact table surrounded by related dimension tables



A **snowflake schema** is an extension of a star, and includes relationships between dimension tables and related sub-dimension tables

ASSIGNMENT: TABLE RELATIONSHIPS



NEW MESSAGE

From: **Dana Modelle** (*Analyst*)

Subject: **Need a favor...**

Hey there,

Ethan shared the data model you've been working on, and we might have an issue...

Last night I left my laptop open, and my cat Dennis somehow got his paws on our model. Now all the relationships are gone!

Could you please rebuild the model, including all three product tables? I owe you one!

-Dana


← Reply

➡ Forward

Key Objectives

1. Delete all existing table relationships
2. Create a star schema by creating relationships between the Sales, Calendar, Customer, Product and Territories tables
3. Connect all three product tables (Product, Subcategory, Category) in a snowflake schema
4. Use the matrix visual to confirm that you can filter Order Quantity values using fields from each dimension table

SOLUTION: TABLE RELATIONSHIPS



NEW MESSAGE

From: **Dana Modelle** (Analyst)

Subject: **Need a favor...**

Hey there,

Ethan shared the data model you've been working on, and we might have an issue...

Last night I left my laptop open, and my cat Dennis somehow got his paws on our model. Now all the relationships are gone!

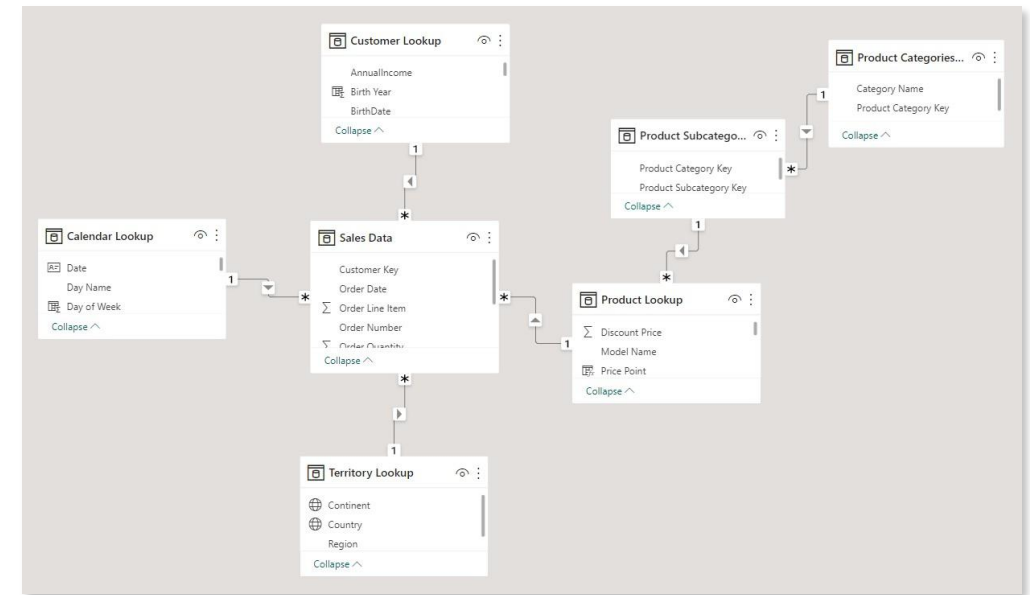
Could you please rebuild the model, including all three product tables? I owe you one!

-Dana

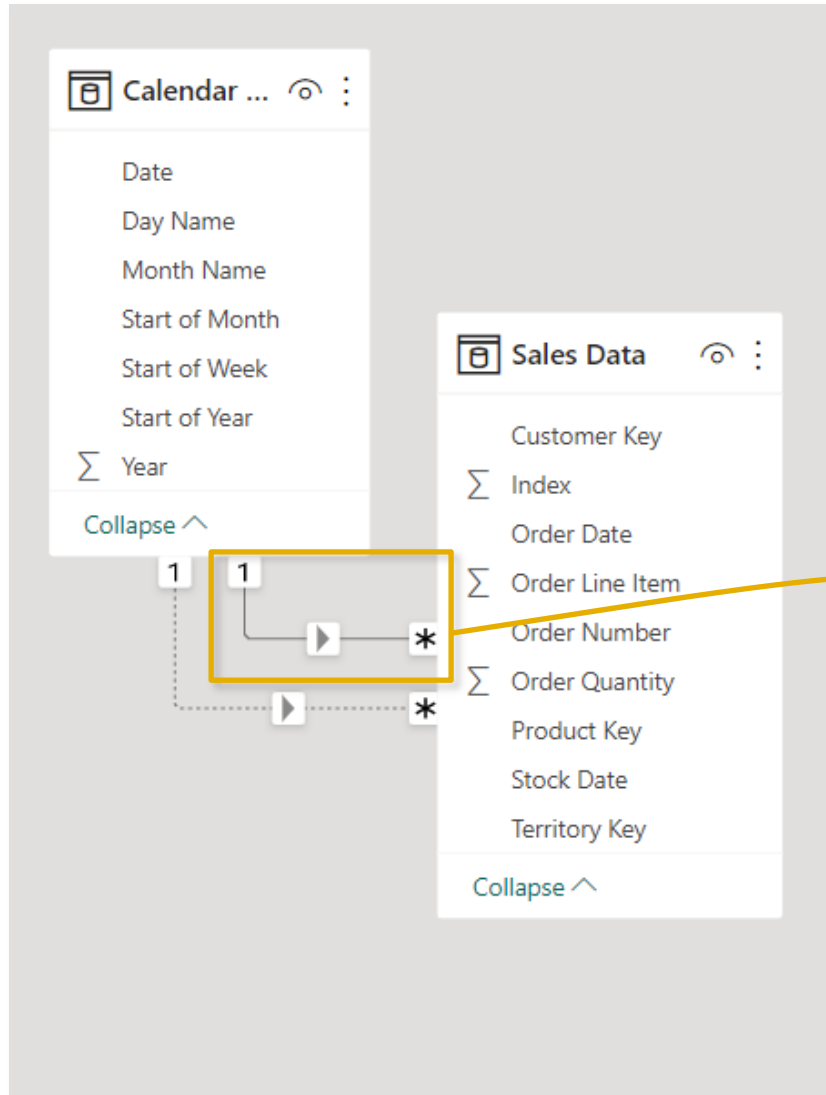
Reply

Forward

Solution Preview



PRO TIP: ACTIVE & INACTIVE RELATIONSHIPS



Edit relationship

Select tables and columns that are related.

Sales Data

| Order Date | Stock Date | Order Number | Product Key | Customer Key | Territory Key | Order Line Item | C |
|------------|------------|--------------|-------------|--------------|---------------|-----------------|---|
| 7/5/2020 | 6/3/2020 | SO46718 | | 360 | 12570 | 9 | 1 |
| 7/7/2020 | 4/22/2020 | SO46736 | | 360 | 12341 | 9 | 1 |
| 7/12/2020 | 5/5/2020 | SO46776 | | 360 | 12356 | 9 | 1 |

Calendar Lookup

| Date | Day Name | Start of Week | Start of Month | Month Name | Start of Year | Year |
|----------|-----------|---------------|----------------|------------|---------------|------|
| 1/1/2020 | Wednesday | 12/29/2019 | 1/1/2020 | January | 1/1/2020 | 2020 |
| 1/2/2020 | Thursday | 12/29/2019 | 1/1/2020 | January | 1/1/2020 | 2020 |
| 1/3/2020 | Friday | 12/29/2019 | 1/1/2020 | January | 1/1/2020 | 2020 |

Cardinality: Many to one (*:1)

Cross filter direction: Single

☒ Make this relationship active

☐ Assume referential integrity

☐ Apply security filter in both directions

OK Cancel

Properties

Relationship

Table: Sales Data

Column: Order Date

Cardinality: Many to one (*:1)

Table: Calendar Lookup

Column: Date

Make this relationship active

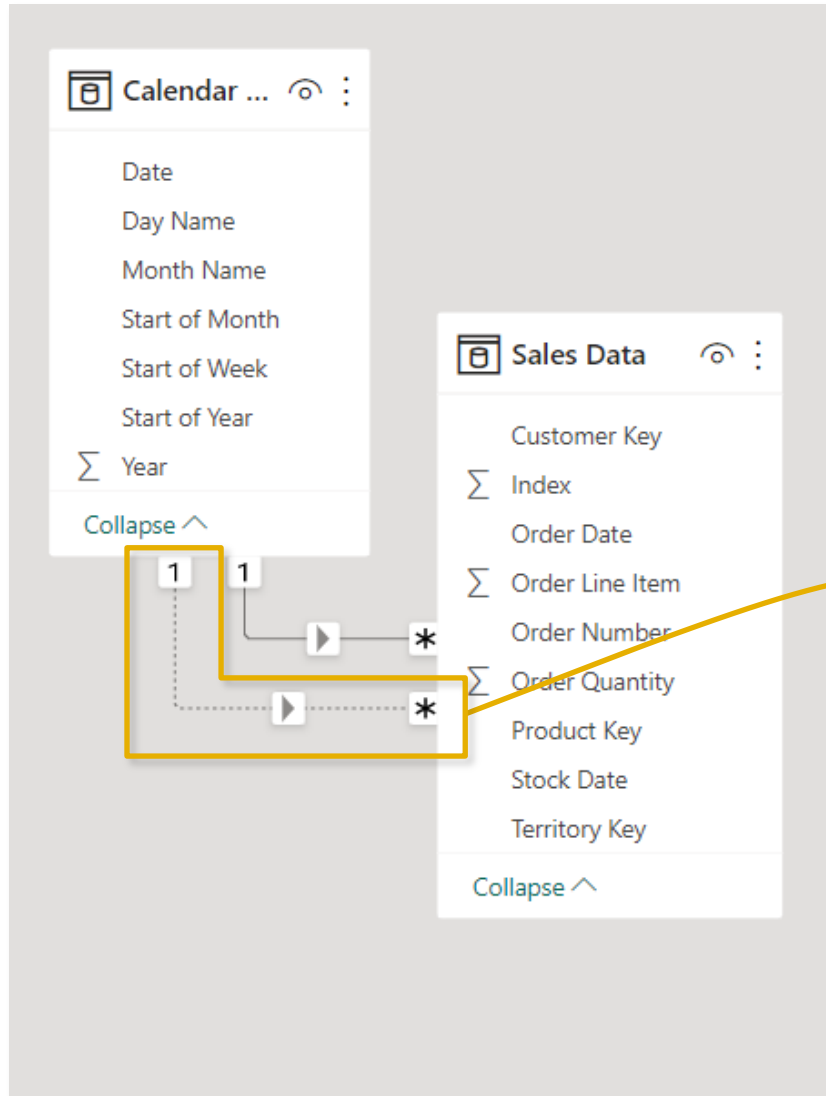
Yes ☒

Cross filter direction: Single

The **Sales Data** table contains two date fields (**Order Date** & **Stock Date**), but there can only be **one active relationship** to the Date key in the Calendar table

You can set relationships to active or inactive from either the **Edit Relationships** dialog box or the **Properties** (you must deactivate one before activating another)

PRO TIP: ACTIVE & INACTIVE RELATIONSHIPS



Edit relationship

Select tables and columns that are related.

Sales Data

| Order Date | Stock Date | Order Number | Product Key | Customer Key | Territory Key | Order Line Item | C |
|------------|------------|--------------|-------------|--------------|---------------|-----------------|---|
| 7/5/2020 | 6/3/2020 | SO46718 | 360 | 12570 | 9 | 1 | |
| 7/7/2020 | 4/22/2020 | SO46736 | 360 | 12341 | 9 | 1 | |
| 7/12/2020 | 5/5/2020 | SO46776 | 360 | 12356 | 9 | 1 | |

Calendar Lookup

| Date | Day Name | Start of Week | Start of Month | Month Name | Start of Year | Year |
|----------|-----------|---------------|----------------|------------|---------------|------|
| 1/1/2020 | Wednesday | 12/29/2019 | 1/1/2020 | January | 1/1/2020 | 2020 |
| 1/2/2020 | Thursday | 12/29/2019 | 1/1/2020 | January | 1/1/2020 | 2020 |
| 1/3/2020 | Friday | 12/29/2019 | 1/1/2020 | January | 1/1/2020 | 2020 |

Cardinality: Many to one (*:1)

Cross filter direction: Single

☐ Make this relationship active

☐ Assume referential integrity

OK Cancel

Properties

Relationship

Table: Sales Data Column: Stock Date

Cardinality: Many to one (*:1)

Table: Calendar Lookup Column: Date

Make this relationship active

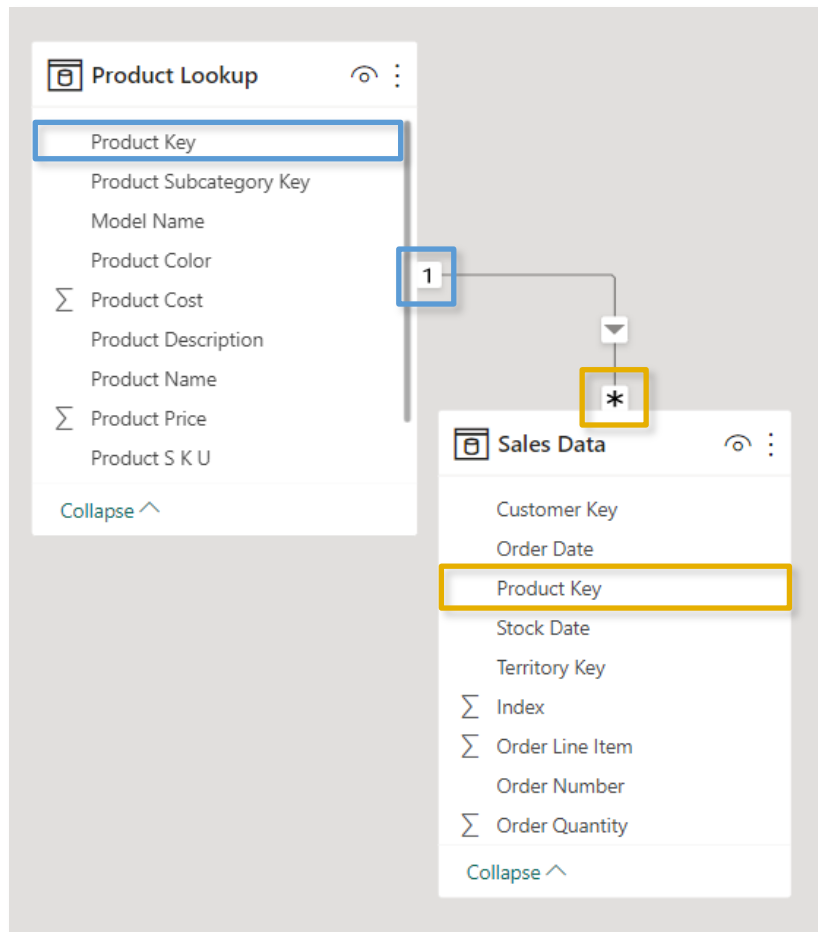
No

Cross filter direction: Single

The **Sales Data** table contains two date fields (**Order Date** & **Stock Date**), but there can only be **one active relationship** to the Date key in the Calendar table

You can set relationships to active or inactive from either the **Edit Relationships** dialog box or the **Properties** (you must deactivate one before activating another)

RELATIONSHIP CARDINALITY



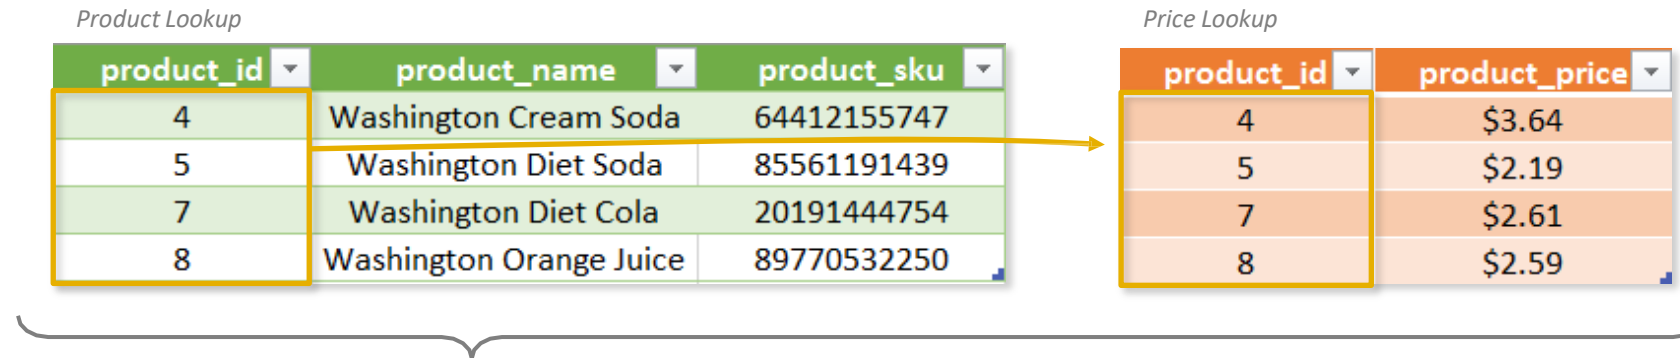
Cardinality refers to the uniqueness of values in a column

- Ideally, all relationships in the data model should follow a **one-to-many** cardinality: **one** instance of each primary key, and **many** instances of each foreign key

*In this example there is only **ONE instance of each Product Key** in the Product table (noted by a “1”), since each row contains **attributes of a single product** (name, SKU, description, price, etc.)*

*There are **MANY instances of each Product Key** in the Sales table (noted by an asterisk *), since there are **multiple sales for each product***

EXAMPLE: ONE-TO-ONE CARDINALITY



- Connecting the two tables above using **product_id** creates a **one-to-one relationship**, since each product ID only appears once in each table
- This isn't necessarily a "bad" relationship, but you can simplify the model by merging the tables into a single, valid dimension table

| product_id | product_name | product_sku | product_price |
|------------|-------------------------|-------------|---------------|
| 4 | Washington Cream Soda | 64412155747 | \$3.64 |
| 5 | Washington Diet Soda | 85561191439 | \$2.19 |
| 7 | Washington Diet Cola | 20191444754 | \$2.61 |
| 8 | Washington Orange Juice | 89770532250 | \$2.59 |

NOTE: this still respects the rules of normalization, since all rows are unique and capture product-specific attributes

EXAMPLE: MANY-TO-MANY CARDINALITY



Product Lookup

| product_id | product_name | product_sku |
|------------|----------------------------|-------------|
| 4 | Washington Cream Soda | 64412155747 |
| 4 | Washington Diet Cream Soda | 81727382373 |
| 5 | Washington Diet Soda | 85561191439 |
| 7 | Washington Diet Cola | 20191444754 |
| 8 | Washington Orange Juice | 89770532250 |

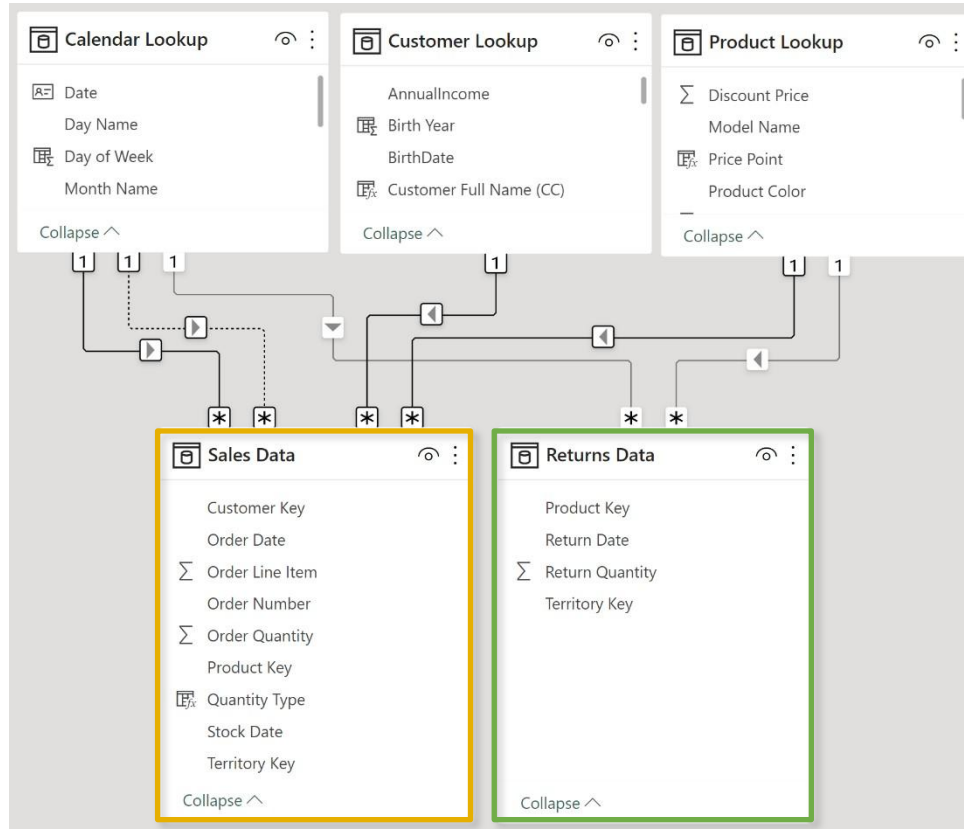
Sales

| date | product_id | transactions |
|----------|------------|--------------|
| 1/1/2017 | 4 | 12 |
| 1/2/2017 | 4 | 9 |
| 1/3/2017 | 4 | 11 |
| 1/1/2017 | 5 | 16 |
| 1/2/2017 | 5 | 19 |
| 1/1/2017 | 7 | 11 |

! This relationship has cardinality Many-Many. This should only be used if it is expected that neither column (product_id and product_id) contains unique values, and that the significantly different behavior of Many-many relationships is understood. [Learn more](#)

- If we try to connect the tables above using **product_id**, we'll get a **many-to-many relationship** warning since there are multiple instances of product_id in both tables
- Even if we force this relationship, how would we know which product was actually sold on each date – **Cream Soda** or **Diet Cream Soda**?

CONNECTING MULTIPLE FACT TABLES



This model contains two fact tables: **Sales Data** and **Returns Data**

- Since there is no primary/foreign key relationship, we can't connect them directly to each other
- But we *can* connect each fact table to related lookups, which allows us to filter both sales and returns data **using fields from any shared lookup tables**
- We can view orders and returns by product since both tables relate to Product Lookup, but we can't view returns by customer since no relationship exists

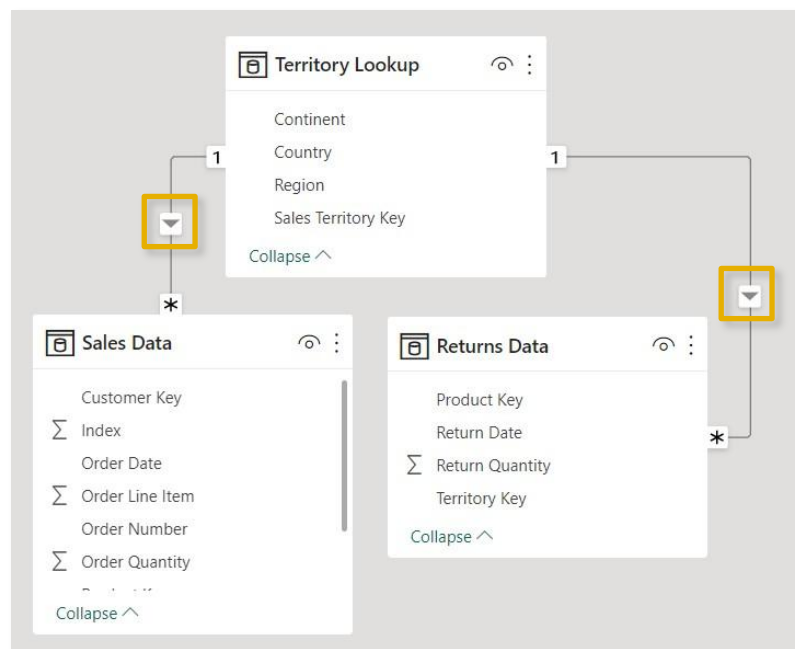


HEY THIS IS IMPORTANT!

Generally speaking, fact tables should **connect through shared dimension tables, not directly to each other**



FILTER CONTEXT & FLOW



Here we have two data tables (**Sales Data** and **Returns Data**), connected to **Territory Lookup**

The arrows show the **filter direction**, and point from the one (1) side of the relationship to the many (*) side

- When you filter a table, that **filter context** is passed to any related “downstream” tables, following the arrow’s direction
- Filter context CANNOT flow “upstream”



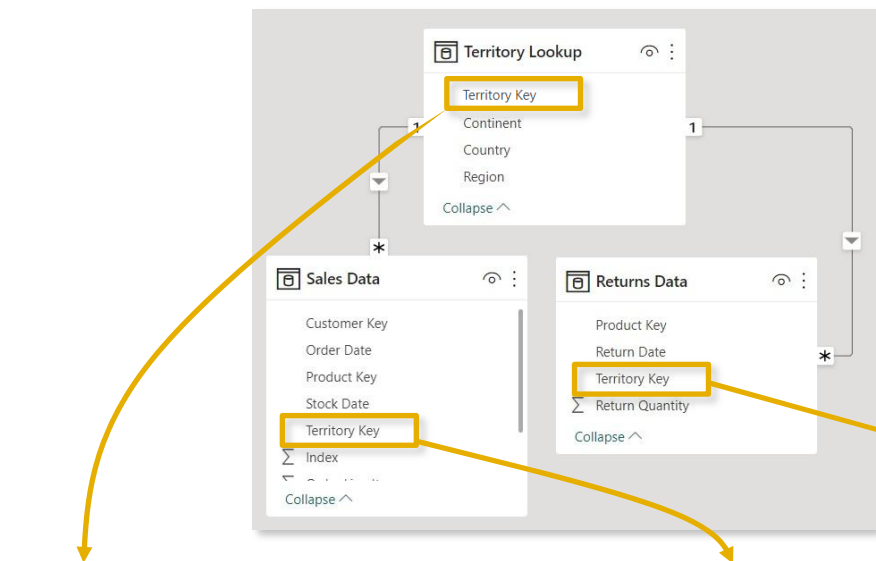
PRO TIP: Arrange lookup tables above fact tables in your model as a visual reminder that **filters always flow downstream**



EXAMPLE: FILTER FLOW

In this model, the only way to filter both **Sales** and **Returns** data by **Territory** is to use the **Territory Key** from the lookup table, which is upstream and related to both fact tables

- Filtering using Territory Key from the **Sales** table yields **incorrect Returns values**, since the filter context can't flow to any other table
- Filtering using Territory Key from the **Returns** table yields **incorrect Sales values**, and is limited to territories that exist in the returns table



| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 12,513 | 270 |
| 2 | 40 | |
| 3 | 30 | |
| 4 | 17,191 | 362 |
| 5 | 17,49 | 1 |
| 6 | 10,894 | 138 |
| 7 | 7,862 | 186 |
| 8 | 7,950 | 163 |
| 9 | 17,951 | 404 |
| 10 | 9,694 | 204 |
| Total | 84,174 | 1,828 |

Filtering by **Territory Lookup**[Territory Key]

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 12,513 | 1,828 |
| 2 | 40 | 1,828 |
| 3 | 30 | 1,828 |
| 4 | 17,191 | 1,828 |
| 5 | 17,49 | 1,828 |
| 6 | 10,894 | 1,828 |
| 7 | 7,862 | 1,828 |
| 8 | 7,950 | 1,828 |
| 9 | 17,951 | 1,828 |
| 10 | 9,694 | 1,828 |
| Total | 84,174 | 1,828 |

Filtering by **Sales Data**[Territory Key]

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 84,174 | 270 |
| 4 | 84,174 | 362 |
| 5 | 84,174 | 1 |
| 6 | 84,174 | 238 |
| 7 | 84,174 | 186 |
| 8 | 84,174 | 163 |
| 9 | 84,174 | 404 |
| 10 | 84,174 | 204 |
| Total | 84,174 | 1,828 |

Filtering by **Returns Data**[Territory Key]



BI-DIRECTIONAL FILTERS

Edit relationship

Select tables and columns that are related.

Sales Data

| Order Date | Stock Date | Order Number | Product Key | Customer Key | Territory Key | Order Quantity |
|-----------------------|---------------------------|--------------|-------------|--------------|---------------|----------------|
| Sunday, July 5, 2020 | Wednesday, June 3, 2020 | SO46718 | 360 | 12570 | 9 | 1 |
| Tuesday, July 7, 2020 | Wednesday, April 22, 2020 | SO46736 | 360 | 12341 | 9 | 1 |
| Sunday, July 12, 2020 | Tuesday, May 5, 2020 | SO46776 | 360 | 12356 | 9 | 1 |

Territory Lookup

| Region | Country | Continent | Sales Territory Key |
|-----------|---------------|---------------|---------------------|
| Northwest | United States | North America | 1 |
| Northeast | United States | North America | 2 |
| Central | United States | North America | 3 |

Cardinality

Many to one (*:1)

☒ Make this relationship active

☐ Assume referential integrity

Cross filter direction

Both

☐ Apply security filter in both directions

OK Cancel

Territory Lookup

- Sales Territory Key
- Continent
- Country
- Region
- Collapse

Sales Data

- Customer Key
- Order Date
- Product Key
- Stock Date
- Territory Key
- Index
- Collapse

Returns Data

- Product Key
- Return Date
- Territory Key
- Return Quantity
- Collapse

Properties

Relationship

Table: Sales Data Column: Territory Key

Cardinality: Many to one (*:1)

Table: Territory Lookup Column: Sales Territory Key

Make this relationship active

Yes ☒

Cross filter direction

Both

Apply security filter in both directions

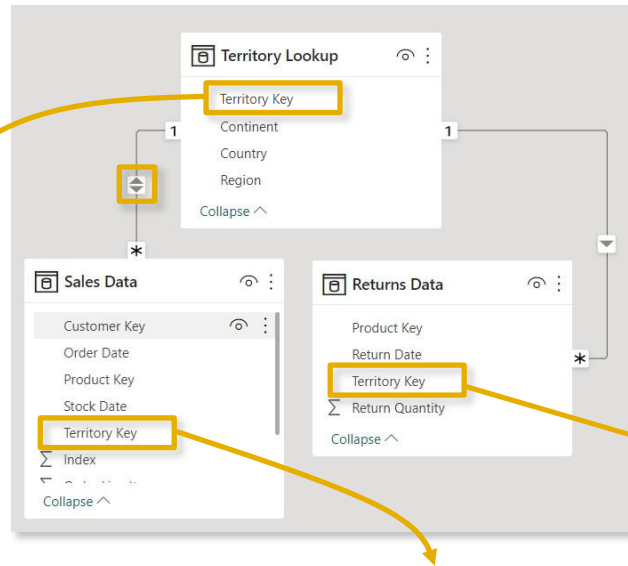
No ☐

Updating the **cross-filter direction** from **Single** to **Both** allows filter context to flow in either direction

- In this example, filters applied to the **Sales** table can pass up to the **Territory Lookup** table, then down to **Returns**



EXAMPLE: BI-DIRECTIONAL FILTERS



With two-way cross-filtering enabled between **Sales** and **Territory**, we now see correct values using **Territory Key** from *either* table

- Filter context can now pass up to the **Territory Lookup** table, then downstream to **Returns**
- However, we still see incorrect values when filtering using Territory Key from the **Returns** table, since the filter context is isolated to that single table

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 12,513 | 270 |
| 2 | 40 | |
| 3 | 30 | |
| 4 | 17,191 | 362 |
| 5 | 49 | 1 |
| 6 | 894 | 238 |
| 7 | 7,862 | 186 |
| 8 | 7,950 | 163 |
| 9 | 17,951 | 404 |
| 10 | 9,694 | 204 |
| Total | 84,174 | 1,828 |

Filtering by **Territory Lookup**[Territory Key]

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 12,513 | 270 |
| 2 | 40 | |
| 3 | 30 | |
| 4 | 17,191 | 362 |
| 5 | 49 | 1 |
| 6 | 894 | 238 |
| 7 | 7,862 | 186 |
| 8 | 7,950 | 163 |
| 9 | 17,951 | 404 |
| 10 | 9,694 | 204 |
| Total | 84,174 | 1,828 |

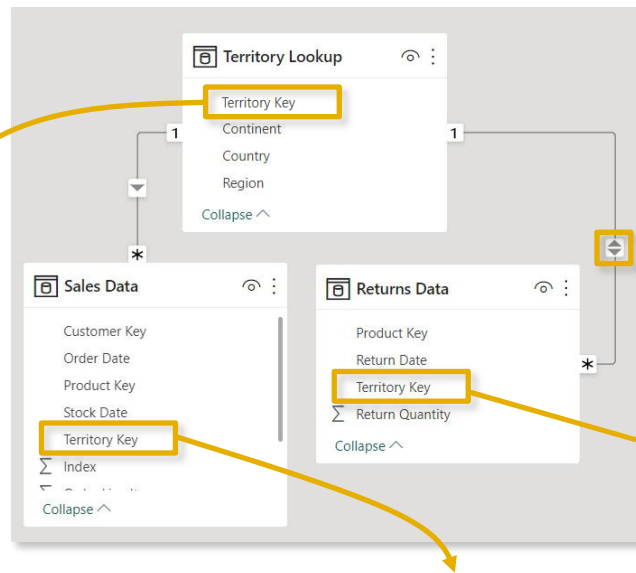
Filtering by **Sales Data**[Territory Key]

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 84,174 | 270 |
| 4 | 84,174 | 362 |
| 5 | 84,174 | 1 |
| 6 | 84,174 | 238 |
| 7 | 84,174 | 186 |
| 8 | 84,174 | 163 |
| 9 | 84,174 | 404 |
| 10 | 84,174 | 204 |
| Total | 84,174 | 1,828 |

Filtering by **Returns Data**[Territory Key]



EXAMPLE: BI-DIRECTIONAL FILTERS



In this case, we've enabled two-way cross-filtering between the **Returns** and **Territory** tables

- As expected, we now see incorrect values when filtering using Territory Key from the **Sales** table, since the filter context is isolated to that single table
- While the values *appear* to be correct when filtering using Territory Key from the **Returns** table, we're **missing sales data** from any territories that didn't appear in the returns table (specifically Territories **2 & 3**)

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 12,513 | 270 |
| 2 | 40 | |
| 3 | 30 | |
| 4 | 17,191 | 362 |
| 5 | 49 | 1 |
| 6 | 10,894 | 238 |
| 7 | 7,862 | 186 |
| 8 | 7,950 | 163 |
| 9 | 17,951 | 404 |
| 10 | 9,694 | 204 |
| Total | 84,174 | 1,828 |

Filtering by **Territory Lookup**[Territory Key]

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 12,513 | 1,828 |
| 2 | 40 | 1,828 |
| 3 | 30 | 1,828 |
| 4 | 17,191 | 1,828 |
| 5 | 49 | 1,828 |
| 6 | 10,894 | 1,828 |
| 7 | 7,862 | 1,828 |
| 8 | 7,950 | 1,828 |
| 9 | 17,951 | 1,828 |
| 10 | 9,694 | 1,828 |
| Total | 84,174 | 1,828 |

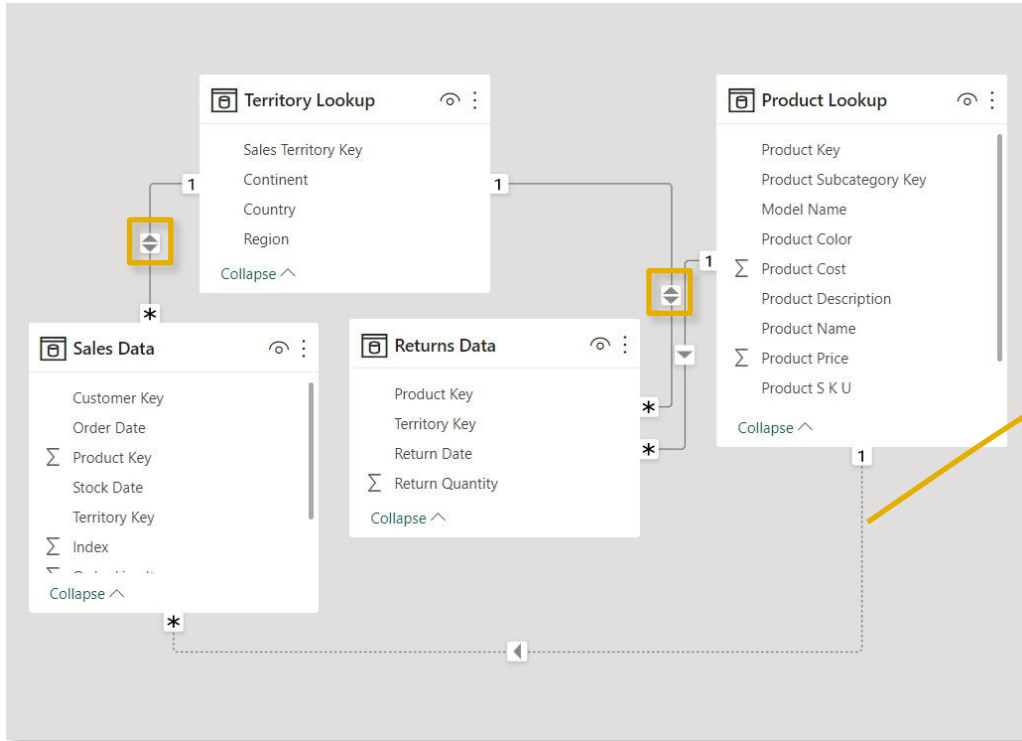
Filtering by **Sales Data**[Territory Key]

| TerritoryKey | OrderQuantity | ReturnQuantity |
|--------------|---------------|----------------|
| 1 | 12,513 | 270 |
| 4 | 17,191 | 362 |
| 5 | 49 | 1 |
| 6 | 10,894 | 238 |
| 7 | 7,862 | 186 |
| 8 | 7,950 | 163 |
| 9 | 17,951 | 404 |
| 10 | 9,694 | 204 |
| Total | 84,174 | 1,828 |

Filtering by **Returns Data**[Territory Key]

Territories 2 & 3 don't exist in the **Returns** table, so they aren't included in the filter context that passes to **Territory Lookup** and **Sales**

AMBIGUITY



Use two-way filters carefully, and **only when necessary**

- Using multiple two-way filters can cause **ambiguity** by introducing multiple filter paths between tables

! You can't create a direct active relationship between **Sales_Data** and **Product_Lookup** because that would introduce ambiguity between tables **Product_Lookup** and **Territory_Lookup**. To make this relationship active, deactivate or delete one of the relationships between **Product_Lookup** and **Territory_Lookup** first.

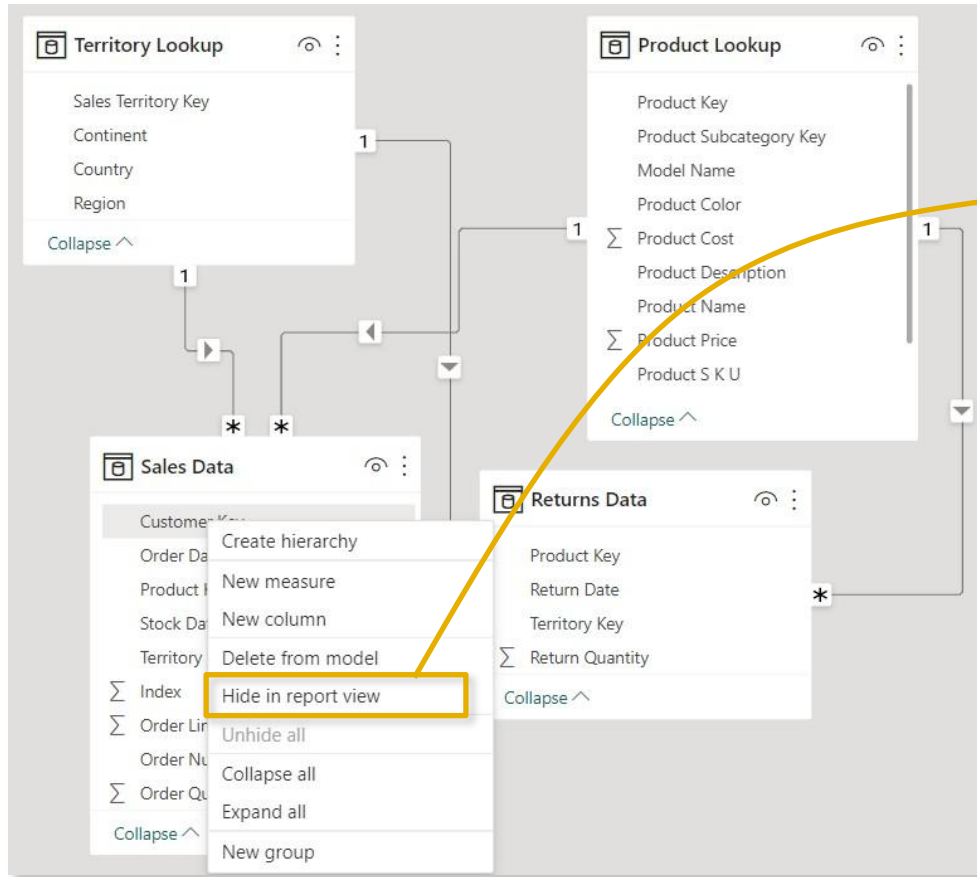
*In this example, filter context from the **Product** table can pass down to **Returns** and up to **Territory Lookup**, which would be filtered based on the Territory Keys passed from the Returns table*

*With an active relationship between **Product** and **Sales** as well, filter context could pass through **either the Sales or Returns table** to reach the **Territory Lookup table**, which could yield conflicting filter context*



PRO TIP: Design your models with **one-way filters** and **1:many cardinality** unless more complex relationships are absolutely necessary

HIDING FIELDS




Hide in Report View makes fields inaccessible from the Report tab, but still available in **Data** and **Model** views

- This can be controlled by right-clicking a field in the Data or Model view, or by selecting “**Is hidden**” in the Properties pane
- This is commonly used to prevent users from filtering using invalid fields, reduce clutter, or to hide irrelevant metrics from view



PRO TIP: Hide the **foreign keys** in fact tables to force users to filter using **primary keys** in dimension tables

ASSIGNMENT: FILTER FLOW



NEW MESSAGE

From: **Dana Modelle** (Analyst)

Subject: **Larry's gone rogue!**

Hey there, we've got another problem.

Larry from Sales just sent me this screenshot. I think he must have downloaded our Power BI model and messed with some relationships, because I KNOW we had sales for product 338.

Can you help diagnose what's going on, and prevent him from doing this again?

-Dana

P.S. Kevin says hi 🐾

← Reply

→ Forward

Key Objectives

1. Replicate Larry's matrix below to diagnose what he must have done to the model*


| Product Key | Sum of Order Quantity | Sum of Return Quantity |
|-------------|-----------------------|------------------------|
| 324 | 72 | 3 |
| 326 | 65 | 3 |
| 328 | 75 | 4 |
| 330 | 51 | 6 |
| 332 | 64 | 2 |
| 334 | 63 | 2 |
| 336 | 50 | 1 |
| 340 | 56 | 1 |
| 342 | 72 | 1 |
| 346 | 24 | 2 |

No sales
for 338!?!
←

- Which product is #338?
 - Why didn't Larry's matrix show any orders?
2. Hide any remaining foreign keys to prevent other users from making the same mistake

**Hint: you may need to temporarily change a relationship to bi-directional*

SOLUTION: FILTER FLOW



NEW MESSAGE

From: **Dana Modelle** (Analyst)

Subject: **Larry's gone rogue!**

Hey there, we've got another problem.

Larry from Sales just sent me this screenshot. I think he must have downloaded our Power BI model and messed with some relationships, because I KNOW we had sales for product 338.

Can you help diagnose what's going on, and prevent him from doing this again?

-Dana

P.S. Kevin says hi 🐾

← Reply

→ Forward

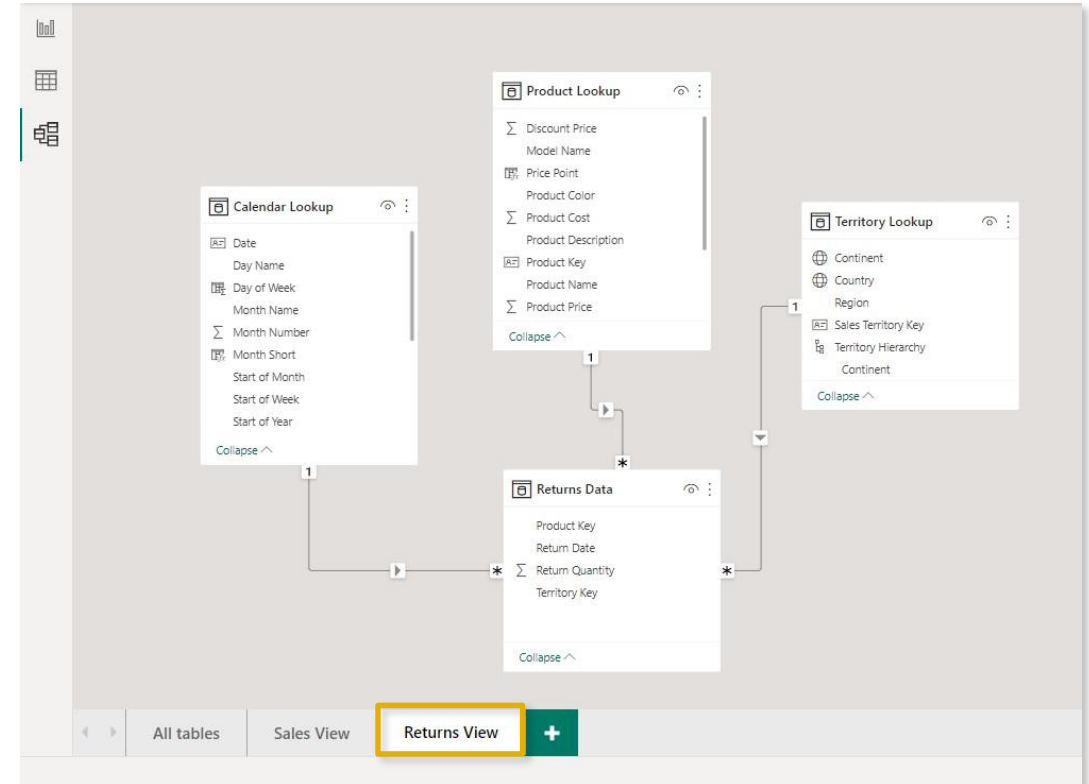
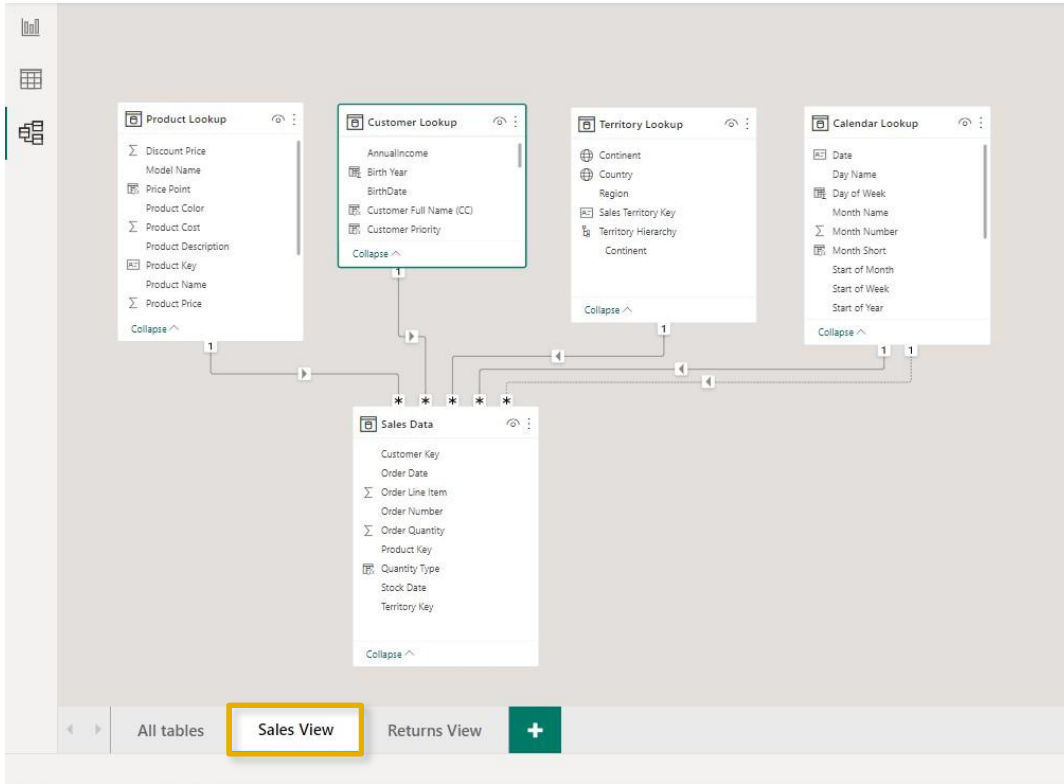
Solution Preview

1. Larry must have changed the relationship between **Returns Data** and **Product Lookup** to **bi-directional**, and filtered his matrix using product_id from the Returns table
 - Road bike (Road-650 Black, 44)
 - Product 338 doesn't exist in the Returns table, so it was excluded when that filter context passed to the Sales table

2.

| Returns Data | |
|-------------------|--|
| Product Key | |
| Return Date | |
| Σ Return Quantity | |
| Σ Territory Key | |

PRO TIP: MODEL LAYOUTS



Model layouts allow you to create custom views to show specific portions of large, complex models

- Here we've created a **Sales View** displaying only tables related to sales, and a **Returns View** displaying only tables related to returns (**Note:** *this doesn't actually create duplicate tables*)

DATA FORMATS & CATEGORIES



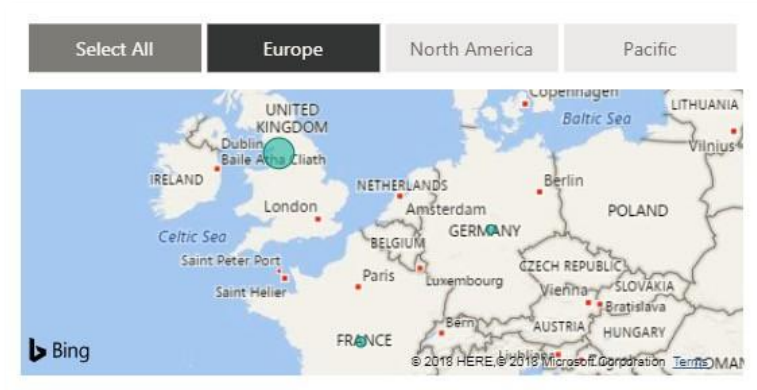
Customize **data formats** from the Column tools menu in the **Data** view or the Properties pane in the **Model** view

The screenshot shows the Power BI interface with the 'Column tools' menu open. The 'Format' dropdown is set to 'Text'. The 'Data category' dropdown is open, showing a list of categories. The 'Country' category is highlighted. The 'Data type' is set to 'Text'. The 'Summarization' is set to 'Don't summarize'. The 'Data category' is set to 'Country'. The 'Sort by' is set to 'column'. The 'Sort' button is visible. The 'Table tools' tab is selected. The 'Column tools' tab is selected. The 'Format' dropdown is set to 'Text'. The 'Data category' dropdown is open, showing a list of categories. The 'Country' category is highlighted. The 'Data type' is set to 'Text'. The 'Summarization' is set to 'Don't summarize'. The 'Data category' is set to 'Country'. The 'Sort by' is set to 'column'. The 'Sort' button is visible. The 'Table tools' tab is selected. The 'Column tools' tab is selected.

| Region | Country | Continent | Sales Territory Key |
|----------------|----------------|---------------|---------------------|
| Northwest | United States | North America | 1 |
| Northeast | United States | North America | 2 |
| Central | United States | North America | 3 |
| Southwest | United States | North America | 4 |
| Southeast | United States | North America | 5 |
| Canada | Canada | North America | 6 |
| France | France | Europe | 7 |
| Germany | Germany | Europe | 8 |
| Australia | Australia | Pacific | 9 |
| United Kingdom | United Kingdom | Europe | 10 |

Assign **data categories** for geospatial fields, URLs or barcodes

- This is commonly used to help Power BI map location-based fields like addresses, countries, cities, coordinates, zip codes, etc.

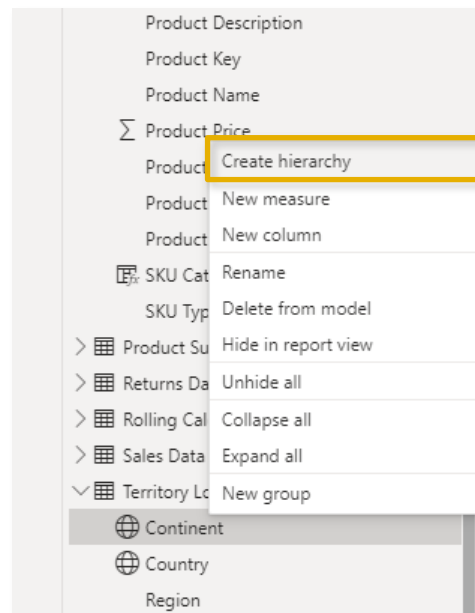




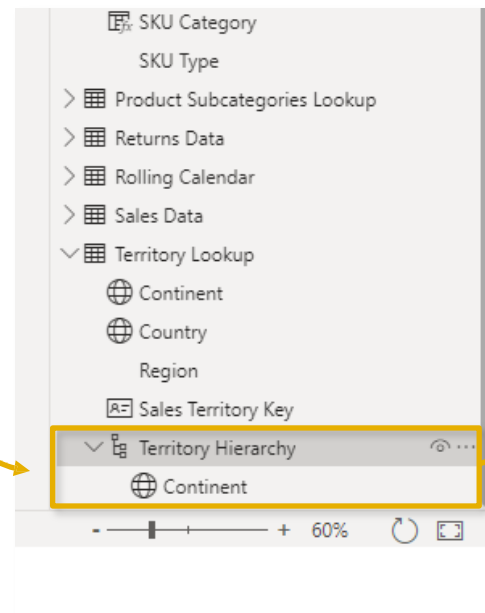
HIERARCHIES

Hierarchies are groups of columns that reflect multiple levels of granularity

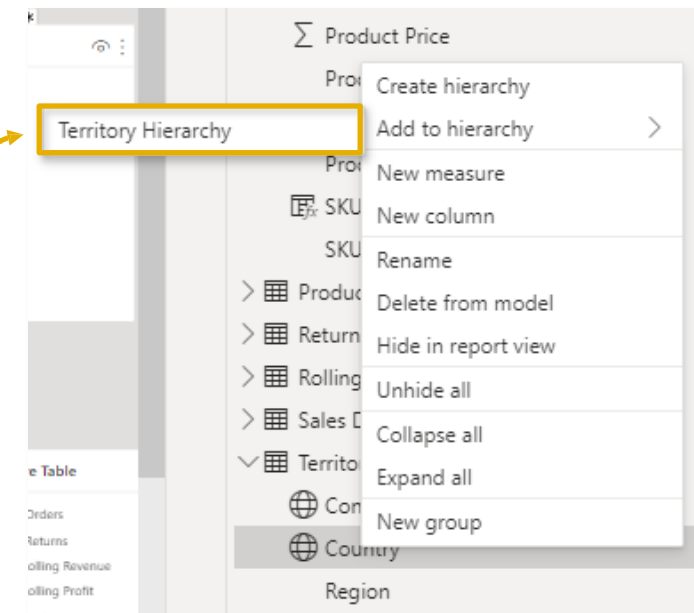
- For example, a **Geography hierarchy** might include **Country**, **State** and **City** fields
- Hierarchies are treated as a **single item** in tables and reports, allowing users to “drill up” and “drill down” through each level



In the **Data** pane, right-click a field and select **Create hierarchy**




This hierarchy contains “Continent”, and is named “**Territory Hierarchy**”



Right-click another field (like “Country”) and select **Add to Hierarchy** (or drag it in!)

ASSIGNMENT: HIERARCHIES



NEW MESSAGE

From: **Dana Modelle** (Analyst)

Subject: **Adding a date hierarchy**

Good morning!

Hoping you can help with a quick request.

Since we'll be doing a lot of time-series analysis, Ethan asked us to add a date hierarchy to the model so that users can quickly view trends at any level of granularity (year, month, day, etc.)

Please get that added before our afternoon call. Thanks!

-Dana

← Reply

➡ Forward


Key Objectives

1. Create a new hierarchy based on the **Start of Year** field, and name it "**Date Hierarchy**"
2. Right-click or drag to add fields until your hierarchy contains the following (in this order):
 - **Start of Year**
 - **Start of Month**
 - **Start of Week**
 - **Date**
3. Add your new hierarchy to the matrix visual (on rows) and practice drilling up and down between each level of granularity

SOLUTION: HIERARCHIES



Solution Preview



NEW MESSAGE

From: **Dana Modelle** (Analyst)

Subject: **Adding a date hierarchy**

Good morning!

Hoping you can help with a quick request.

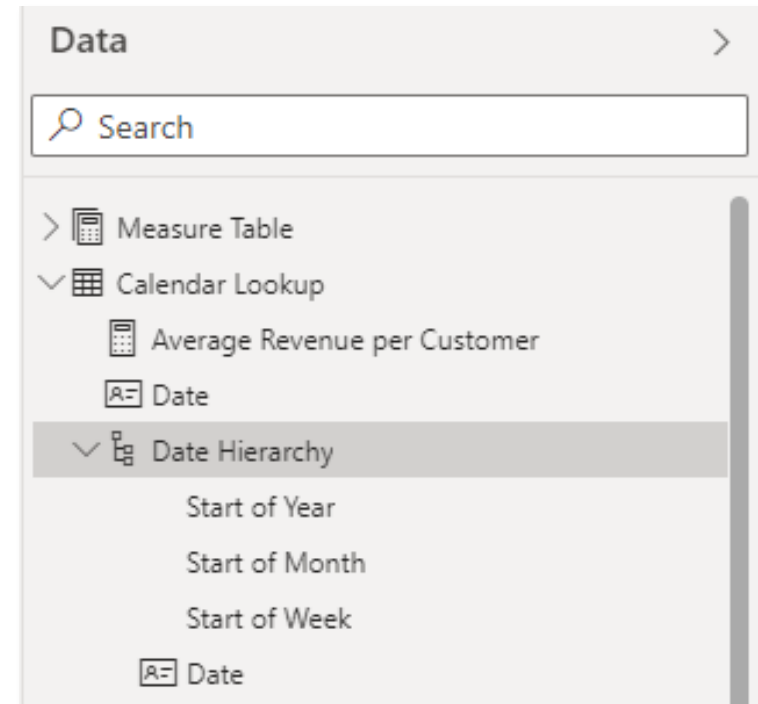
Since we'll be doing a lot of time-series analysis, Ethan asked us to add a date hierarchy to the model so that users can quickly view trends at any level of granularity (year, month, day, etc.)

Please get that added before our afternoon call. Thanks!

-Dana

← Reply

→ Forward



DATA MODEL BEST PRACTICES



- ★ Focus on building a normalized model from the start
 - *Leverage relationships and make sure that each table serves a clear, distinct purpose*
- ★ Organize dimension tables above data tables in your model
 - *This serves as a visual reminder that filters always flow “downstream”*
- ★ Avoid complex relationships unless absolutely necessary
 - *Aim to use 1-to-many table relationships and one-way filters whenever possible*
- ★ Hide fields from report view to prevent invalid filter context
 - *This forces report users to filter using primary keys from dimension tables*